

1. Fully associative (FA) caches
  - a. Any block of main memory can go into any line
  - b.  $W = LIC, S = 1$ 
    - i. Entire cache is a whole set
    - ii. Any line of cache can go into any way of that set
    - iii. Number of set bits are always 0
  - c. Won't have conflicts like with a DM cache
    - i. Any line can go anywhere
  - d. However, extremely expensive to implement in terms of both power and money
    - i. Must check all tags at once to see if data is in cache
  - e. Example
    - i. Same cache parameters as before, except now a FA cache
    - ii. 8-byte FA cache with line size of 2, and 4-bit address
    - iii.  $LIC = C / LS = 8 / 2 = 4$
    - iv. For a FA cache,  $S = 1$

Tag	Set	Offset	Address Bits
$4 - 1 - 0 = 3$ bits	$\log_2 S = \log_2 1 = 0$ bits	$\log_2 LS = \log_2 2 = 1$ bit	= 4 bits

2. FA cache mapping example (*not covered in lecture*)
  - a. Use the same cache specifications and address format we calculated in the previous problem
  - b. Use same addresses and data from DM cache example previously

Address	Data	Address	Data
0110	0x1B	0111	0x59
1000	0xFE	1001	0x3D
1010	0x25	1011	0x98
1110	0x0C	1111	0x3A

- c. Load addresses into cache in same fashion
  - a. However, this time there's no set bits
  - b. Any line of cache can go anywhere, as we see in the next table

Memory Access	Tag	Offset	Access Type	Data	Hit or Miss
0110	011	0	Read		Miss
1000	100	0	Read		Miss
1110	111	0	Read		Miss
1010	101	0	Write	0xBE	Miss
1011	101	1	Write	0xEF	Hit

- d. Place values from first three accesses into table below
- First three accesses fill up first three previously invalid lines of cache
  - Each one is a miss, but they bring in the entire line
  - No conflict between 0110 and 1110 this time
  - Dirty bits are all 0 for same reason as before

Line in Cache	Tag	Byte 0	Byte 1	Valid	Dirty
00	011	1B	59	0 1	0 0
01	100	FE	3D	0 1	0 0
10	111	0C	3A	0 1	0 0
11				0	0

- e. Now we introduce writes
- Write changes the value in the cache
  - The first write misses, so we must bring the original values in from memory
    - Known as *write-allocate* policy
    - Once we bring it in, *then* we modify the 0<sup>th</sup> byte
    - Set dirty bit for this line since we modified the value from memory
  - Second write hits, so we don't need to bring the line in
    - Change the 1<sup>st</sup> byte to the new value
    - Dirty bit stays 1 here

Line in Cache	Tag	Byte 0	Byte 1	Valid	Dirty
00	011	1B	59	0 1	0 0
01	100	FE	3D	0 1	0 0
10	111	0C	3A	0 1	0 0
11	101	25 BE	98 EF	0 1	0 1

3. Set associative (SA) caches
- Compromise between DM and FA
  - N-way SA cache means there's N lines in each set
    - Thus, there's N different places a line can go into
    - $W = N(-\text{way})$
  - Advantages and disadvantages of both DM and FA
    - More options to place lines, so less conflicts than a DM cache
    - Will still have conflicts compared to a FA cache
    - More expensive in power and cost than a DM cache
    - Cheaper and less power than a FA cache
  - Example
    - Same cache parameters as before, except now a 2-way SA cache
    - 8-byte 2-way SA cache with line size of 2, and 4-bit address
    - $LIC = C / LS = 8 / 2 = 4$
    - For a SA cache,  $S = LIC / W = 4 / 2 = 2$

Tag	Set	Offset	Address Bits
$4 - 2 - 0 = 2 \text{ bits}$	$\log_2 S = \log_2 2 = 1 \text{ bit}$	$\log_2 LS = \log_2 2 = 1 \text{ bit}$	= 4 bits

4. SA cache mapping example (*not covered in lecture*)
  - a. Use the same cache specifications and address format we calculated in the previous problem
  - b. Use same addresses and data from previous two examples

Address	Data	Address	Data
0110	0x1B	0111	0x59
1000	0xFE	1001	0x3D
1010	0x25	1011	0x98
1110	0x0C	1111	0x3A

- c. Load addresses into cache into same fashion
  - i. One set bit this time

Memory Access	Tag	Set	Offset	Access Type	Data	Hit or Miss
0110	01	1	0	Read		Miss
1000	10	0	0	Read		Miss
1110	11	1	0	Read		Miss
1010	10	1	0	Read		Miss
1110	11	1	0	Write	DE	Hit

- d. Place values from accesses into table below
  - i. Table stops right before the fourth read
  - ii. How do we know which line to evict in set 1?

Set	Line in Cache	Tag	Byte 0	Byte 1	Valid	Dirty
0	0	10	FE	3D	0 1	0 0
	1				0	0
1	2	01	1B	59	0 1	0 0
	3	11	0C	3A	0 1	0 0

5. Cache replacement algorithms
  - b. How do we evict lines from a given set for non-DM caches?
    - i. Why do we not need a replacement policy for DM caches?
  - c. Least recently used (LRU)
    - i. Whichever line that has gone the longest without being touched is evicted
    - ii. 2-way cache example: when we touch a line, set its *use* bit and clear the other line's bit
  - d. First in, first out (FIFO)
    - i. Evict whichever line has been in the cache the longest
    - ii. Use a queue to implement
  - e. Least frequently used (LFU)
    - i. Add a counter to each line
    - ii. Evict whichever line has the fewest accesses
  - f. Random
    - i. Works almost as well as any of the others

6. Finishing the SA cache mapping example (*not covered in lecture*)
- Will assume LRU for our previous example
    - Have added correct use bit history below
    - Any access sets that line's use bit to 1 and the other line's bit to 0
      - Set 1's use bits switch twice with the second and third accesses

Set	Line in Cache	Tag	Byte 0	Byte 1	Valid	Dirty	Use
0	0	10	FE	3D	01	00	01
	1				0	0	00
1	2	01	1B	59	01	00	010
	3	11	0C	3A	01	00	001

- Fourth read needs to evict a line in set 1
  - Choose one with use bit 0
    - Was the first line we brought in, so was "least recently used"
  - After bringing in new line, set its use bit to 1 since we just accessed it
  - Change other line's use bit to 0

Set	Line in Cache	Tag	Byte 0	Byte 1	Valid	Dirty	Use
0	0	10	FE	3D	01	00	01
	1				0	0	00
1	2	<del>01</del>	<del>1B</del>	59	<del>01</del>	00	<del>010</del>
		10	25	98	1	0	1
	3	11	0C	3A	01	00	0010

- Fifth write is a hit
  - Set dirty bit appropriately
  - Use bit switches again

Set	Line in Cache	Tag	Byte 0	Byte 1	Valid	Dirty	Use
0	0	10	FE	3D	01	00	01
	1				0	0	00
1	2	<del>01</del>	<del>1B</del>	59	<del>01</del>	00	<del>010</del>
		10	25	98	1	0	10
	3	11	<del>0C</del> DE	3A	01	001	00101

7. Cache write policies

- a. If we change values that reside in main memory, we make changes in cache first
- b. What happens when we need to evict a line?
- c. Write-back
  - i. Only update RAM if evicted line's dirty bit is set
  - ii. Dirty bit set if CPU had to update value of line in cache
  - iii. Does not require writing to RAM nearly as often
    - 1. If we don't write to a line, then evict it, don't need to update main memory
  - iv. However, RAM value doesn't always match the updated cache value anymore
    - 1. Other devices that need the most up-to-date value must go through cache now
- d. Write-through
  - i. All write operations are made to main memory as well as cache
  - ii. Every time a value is changed, must update RAM as well
  - iii. Requires writing to RAM extremely often
  - iv. Write-through does *not* use a dirty bit